

# Machine learning for ecology with R


Brett Melbourne  
Associate Professor, EBIO  
[brett.melbourne@colorado.edu](mailto:brett.melbourne@colorado.edu)  
Pronouns: he, him, his

# GitHub

- Code for this presentation
  - [github.com/melbourne-lab/ml4e-nutshell](https://github.com/melbourne-lab/ml4e-nutshell)
- Machine learning for ecology graduate course
  - [github.com/EBIO5460Spring2022](https://github.com/EBIO5460Spring2022)

# What is machine learning?


- Using generic algorithms to **predict** outputs  $y$  from inputs  $x$
- Emphasis: predictive skill



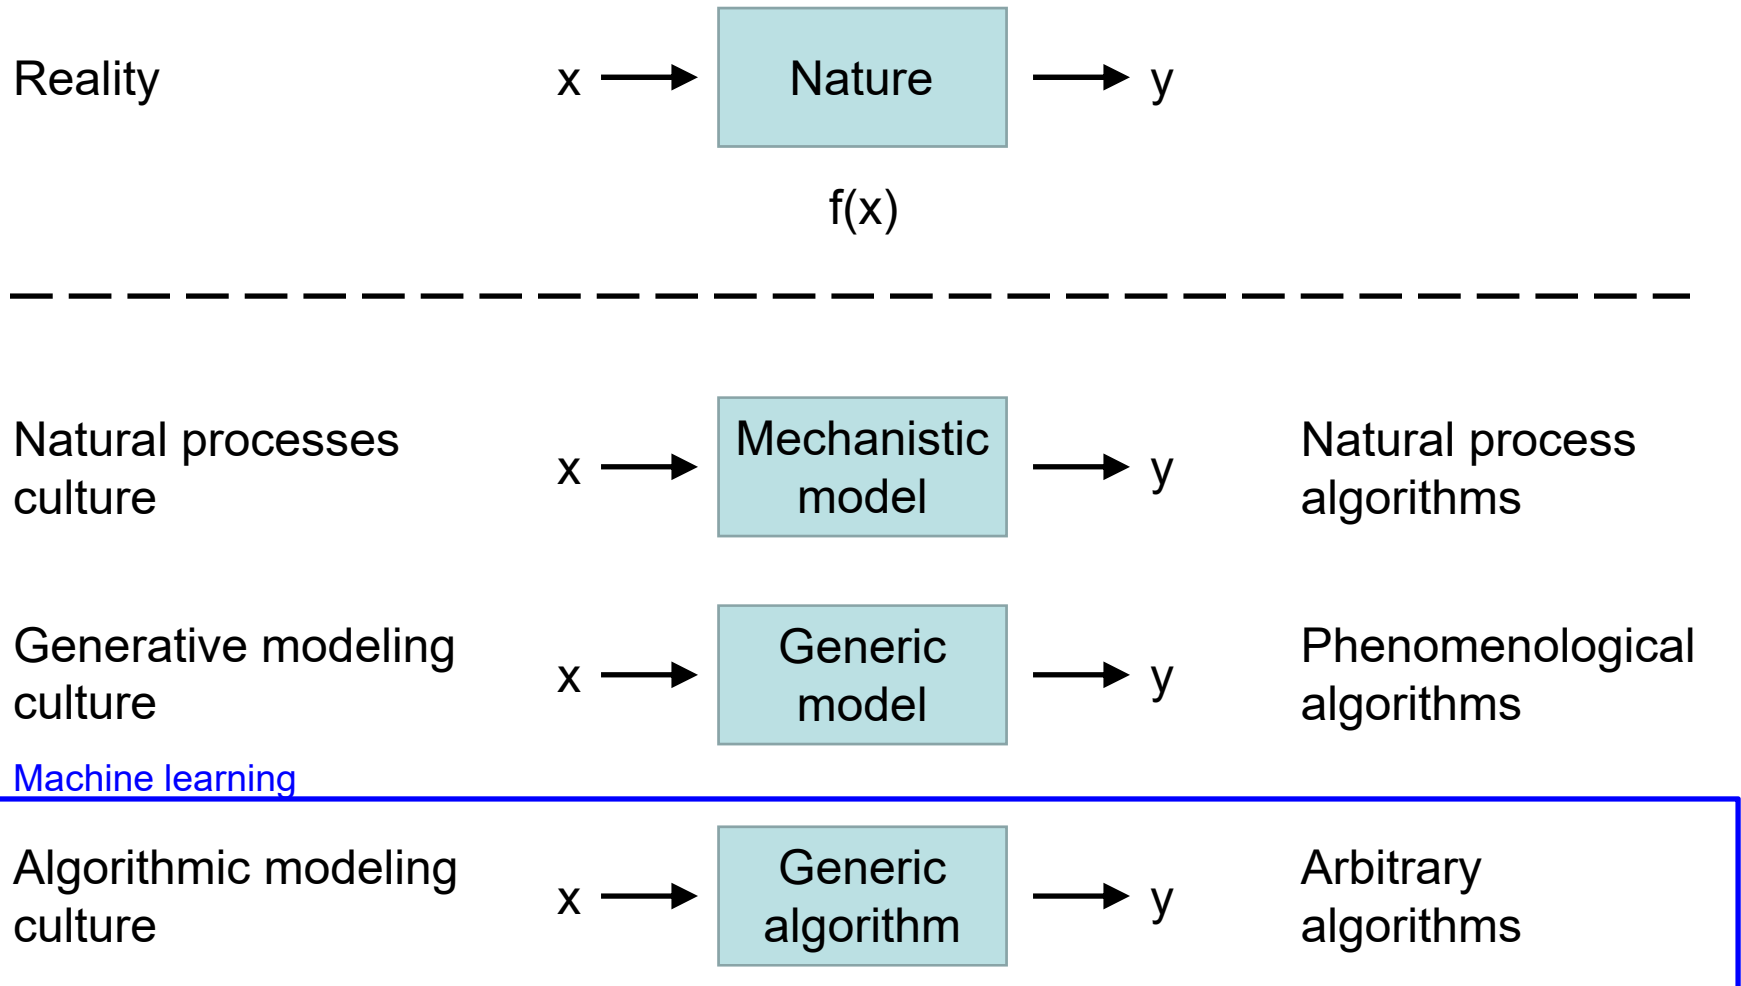
# Examples in ecology

- Species distribution models (SDMs)
  - predicting the spatial distribution of a species from explanatory variables
- Counting the number of penguins in Antarctica from satellite imagery
- Identifying species in camera trap images
- Identifying bird species from audio recordings

# Trying to learn a function $f$



# Trying to learn a function $f$



$f$  can mean different things in different data science cultures

# Goal of prediction

Use data to find a function  $\hat{f}$  that has good predictive performance given  $X$

That is,  $\hat{f}$  is accurate on new observations

# Goal of machine learning

To predict accurately!


- Species distribution
  - map
  - predict accurately for places we won't visit
- Climate change forecast
  - predict accurately for the future
- Antelopes in camera trap images
  - hand over the identification task to a machine so we don't have to look at images!
  - predict accurately for images that we'll never look at




# Predictive skill

Basic idea: out-of-sample accuracy

$\hat{f}$  fitted on training data



$\hat{f}$  predicting new data



e.g. mean square error (MSE)  $\frac{1}{n} \sum_{i=1}^n e_i^2$

# Basic full ML setup


## Overall algorithm:

1. Create a **model algorithm** for  $\hat{f}(x)$
2. Use a **training algorithm** to find parameter values of  $\hat{f}(x)$
3. Use an **inference algorithm** to compare predictive skill among models (model families, tuning parameters,  $x$  sets, etc).


# Inference algorithm

Basic idea: out-of-sample validation

Fit model to training dataset



Test model on validation dataset




e.g. mean square error (MSE)

# k-fold cross validation (CV)

Divide dataset into k parts (preferably randomly)

test (validation)  
data

training  
data



repeat with  
next test subset


... repeat with each test subset

# Regression & classification

- Regression:
  - numerical response variable
  - predict a numerical value given  $x$
  - e.g. number of species given latitude
- Classification:
  - categorical response variable
  - predict the category given  $x$
  - e.g. is it a bird, deer, tree, or mountain lion?
  - e.g. is it dead or alive?; present or absent?

# Ants data

```
> head(ants)
  site habitat latitude elevation richness
1  TPB  forest   41.97      389         6
2  HBC  forest   42.00         8        16
3  CKB  forest   42.03      152        18
4  SKP  forest   42.05         1        17
5   CB  forest   42.05      210         9
6  RP  forest   42.17         78        15
```



# Basic full ML setup


- Polynomial example, 3 algorithms:
  - **model**: flexible function  $\hat{f}(x)$ ;  
polynomial linear model

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \dots + \beta_m x^m \quad m=\text{order}$$




# Ants (forest habitat)


Polynomial order 1




Polynomial order 2



Polynomial order 3



Polynomial order 4



# Basic full ML setup

- Polynomial example, 3 algorithms:

- **model**: flexible function  $\hat{f}(x)$ ;  
polynomial linear model

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \dots + \beta_m x^m \quad m=\text{order}$$

- **training**: optimize least squares objective  
function

- minimize  $SSQ = \sum_{i=1}^n (y_i - \hat{y}_i)^2$  for training  
data

```
lm(richness ~ poly(latitude, order), data=forest_ants)
```

# Basic full ML setup

- Polynomial example, 3 algorithms:

- **model**: flexible function  $\hat{f}(x)$ ;  
polynomial linear model

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \dots + \beta_m x^m \quad m=\text{order}$$

- **training**: optimize least squares objective function

- minimize  $SSQ = \sum_{i=1}^n (y_i - \hat{y}_i)^2$  for training data


- **inference**: tuning parameter (order of poly);  
k-fold cross validation

# Inference algorithm

## k-fold cross validation

```
cv_ants <- function(k, order) {  
  forest_ants$fold <- random_folds(nrow(forest_ants), k)  
  e <- rep(NA, k)  
  for ( i in 1:k ) {  
    test_data <- forest_ants %>% filter(fold == i)  
    train_data <- forest_ants %>% filter(fold != i)  
    poly_trained <- lm(richness ~ poly(latitude, order), data=train_data)  
    pred_richness <- predict(poly_trained, newdata=test_data)  
    e[i] <- mean((test_data$richness - pred_richness) ^ 2)  
  }  
  cv_error <- mean(e)  
  return(cv_error)  
}
```


## Leave one out CV (k-fold CV, $k=n$ )




# Basic ML algorithms


- Ensemble methods
  - Bagging
  - Random forest
  - Boosting
- Neural networks

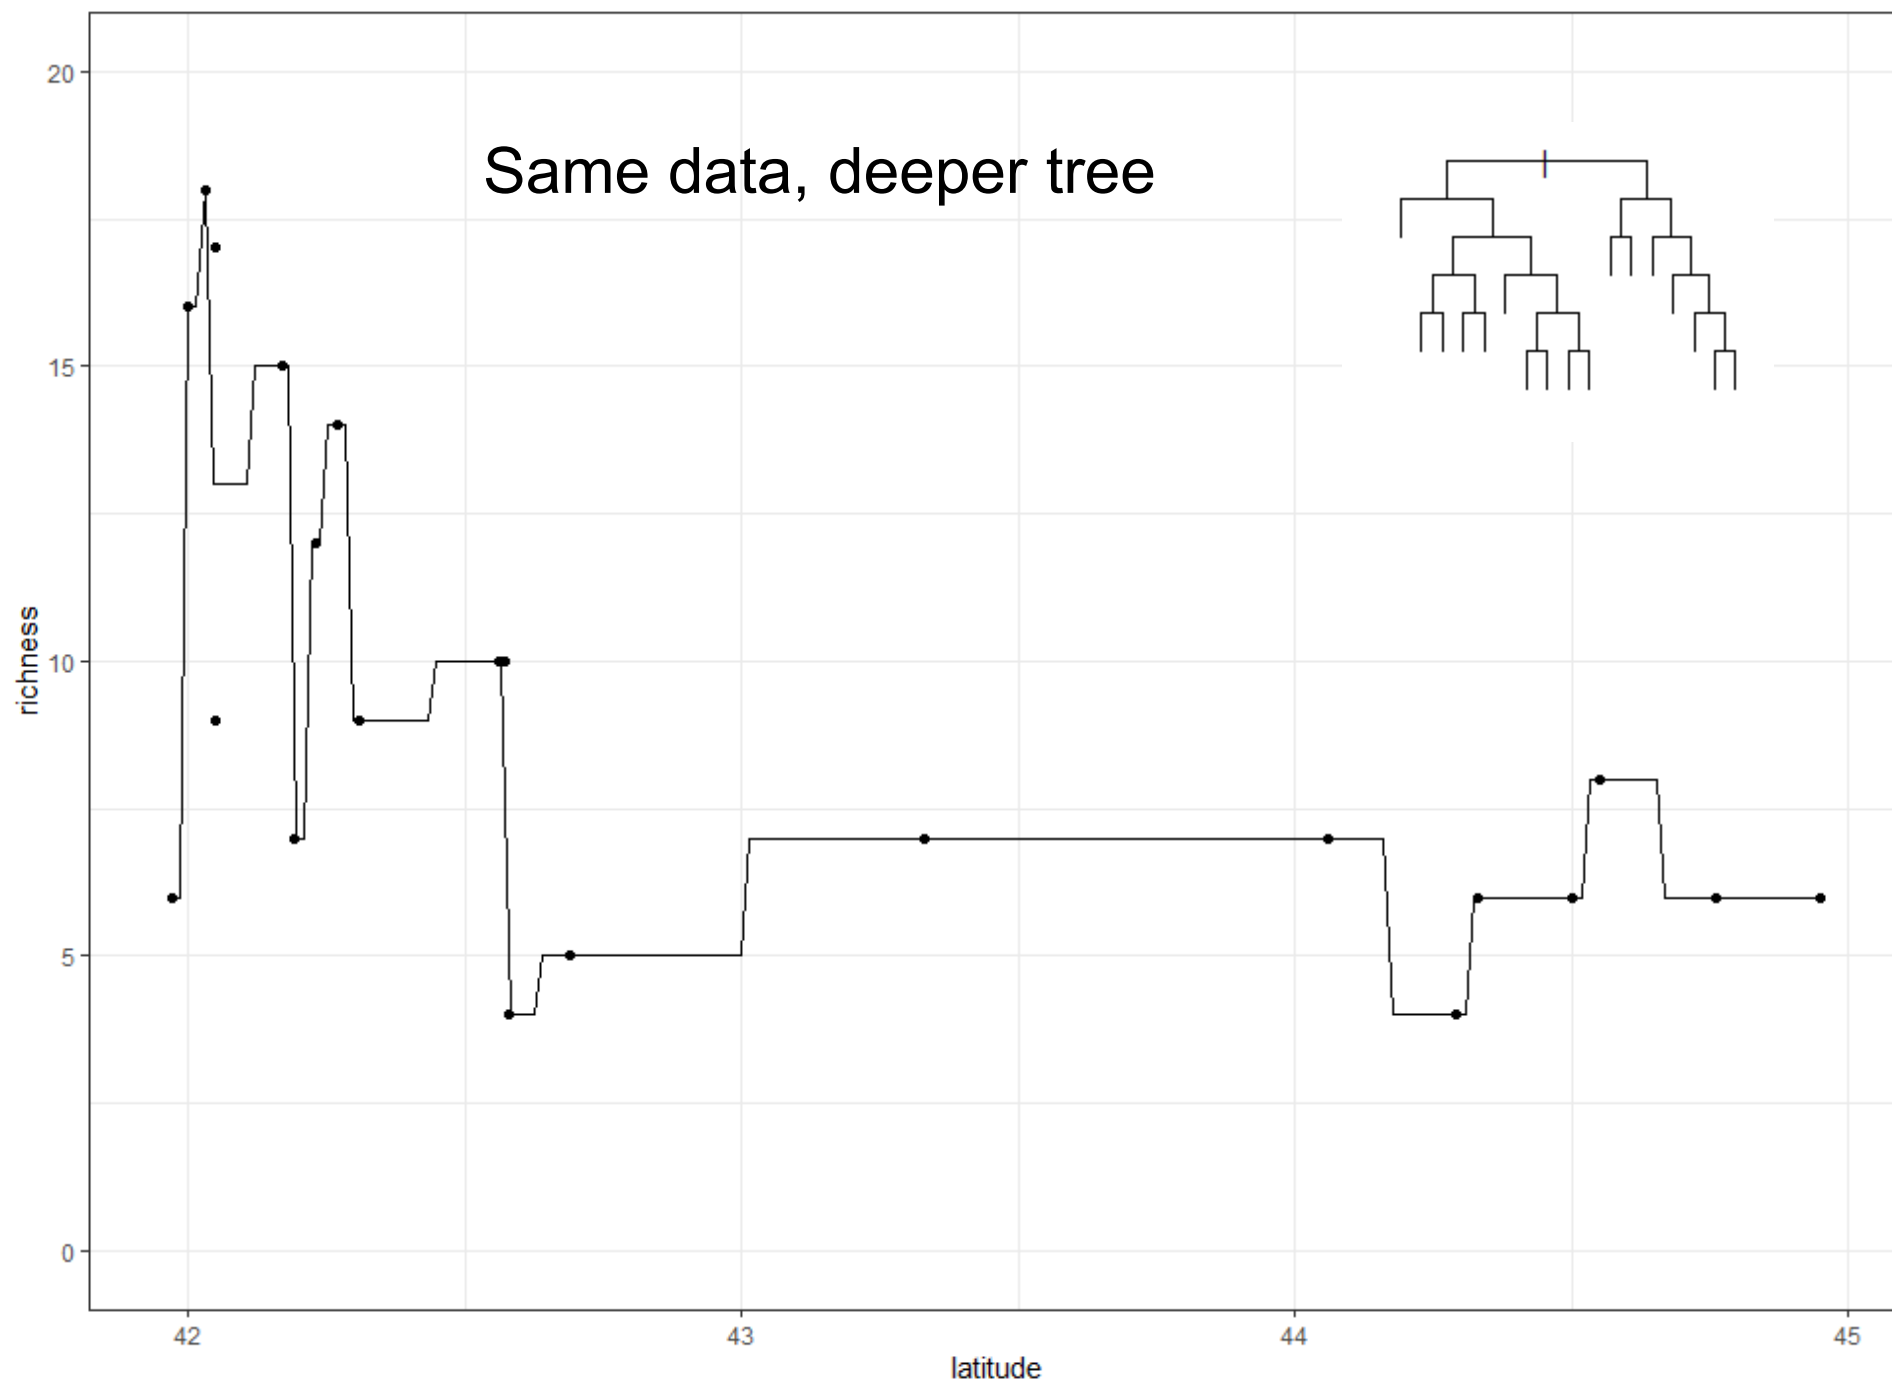
# Regression tree base model






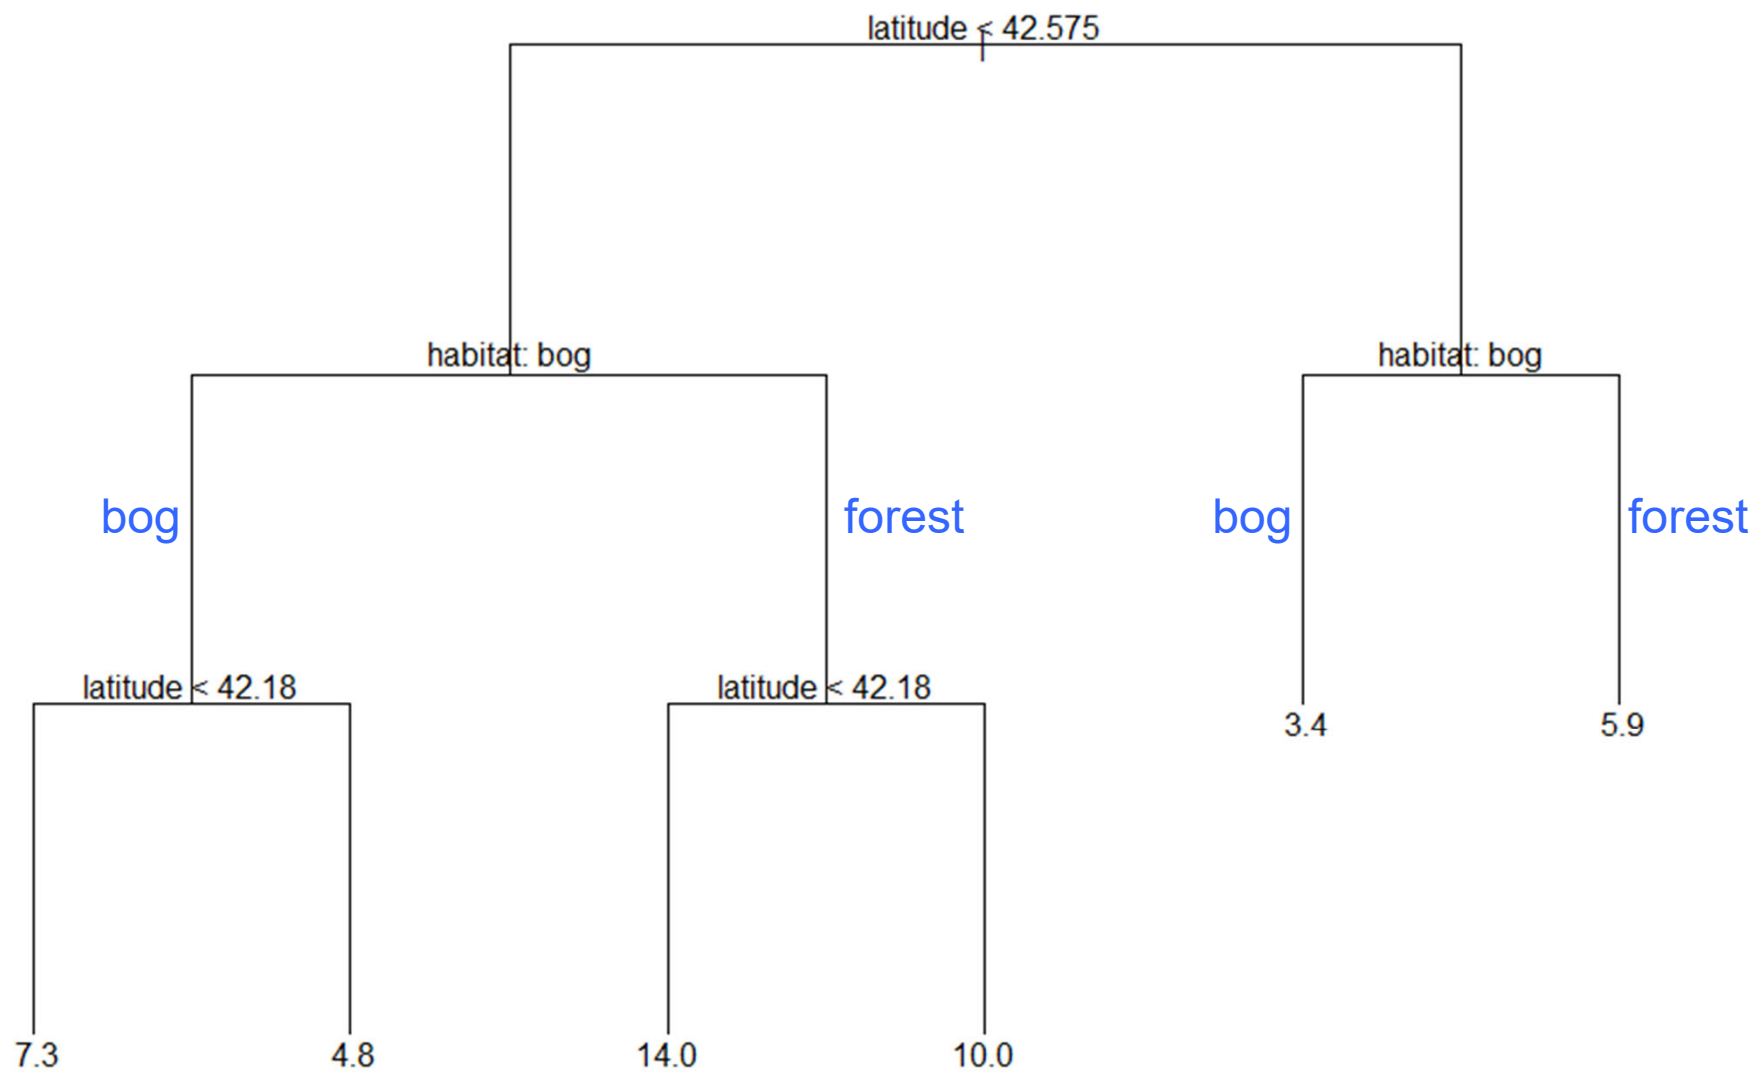






## Nonlinearity or “interaction” of habitat and latitude





# Ensemble methods

- Train many models
- Average the models to predict
- Averaging reduces variance

e.g.  $\text{Var}(\bar{y}) = \frac{\sigma_y^2}{n}$

# Bagging

- Bootstrap
  - form new datasets by resampling from the data
- Aggregate
  - average over bootstrapped model fits

# Bagging algorithm

for many repetitions

- resample the data with replacement


- train the base model

- record prediction

final prediction = mean of predictions

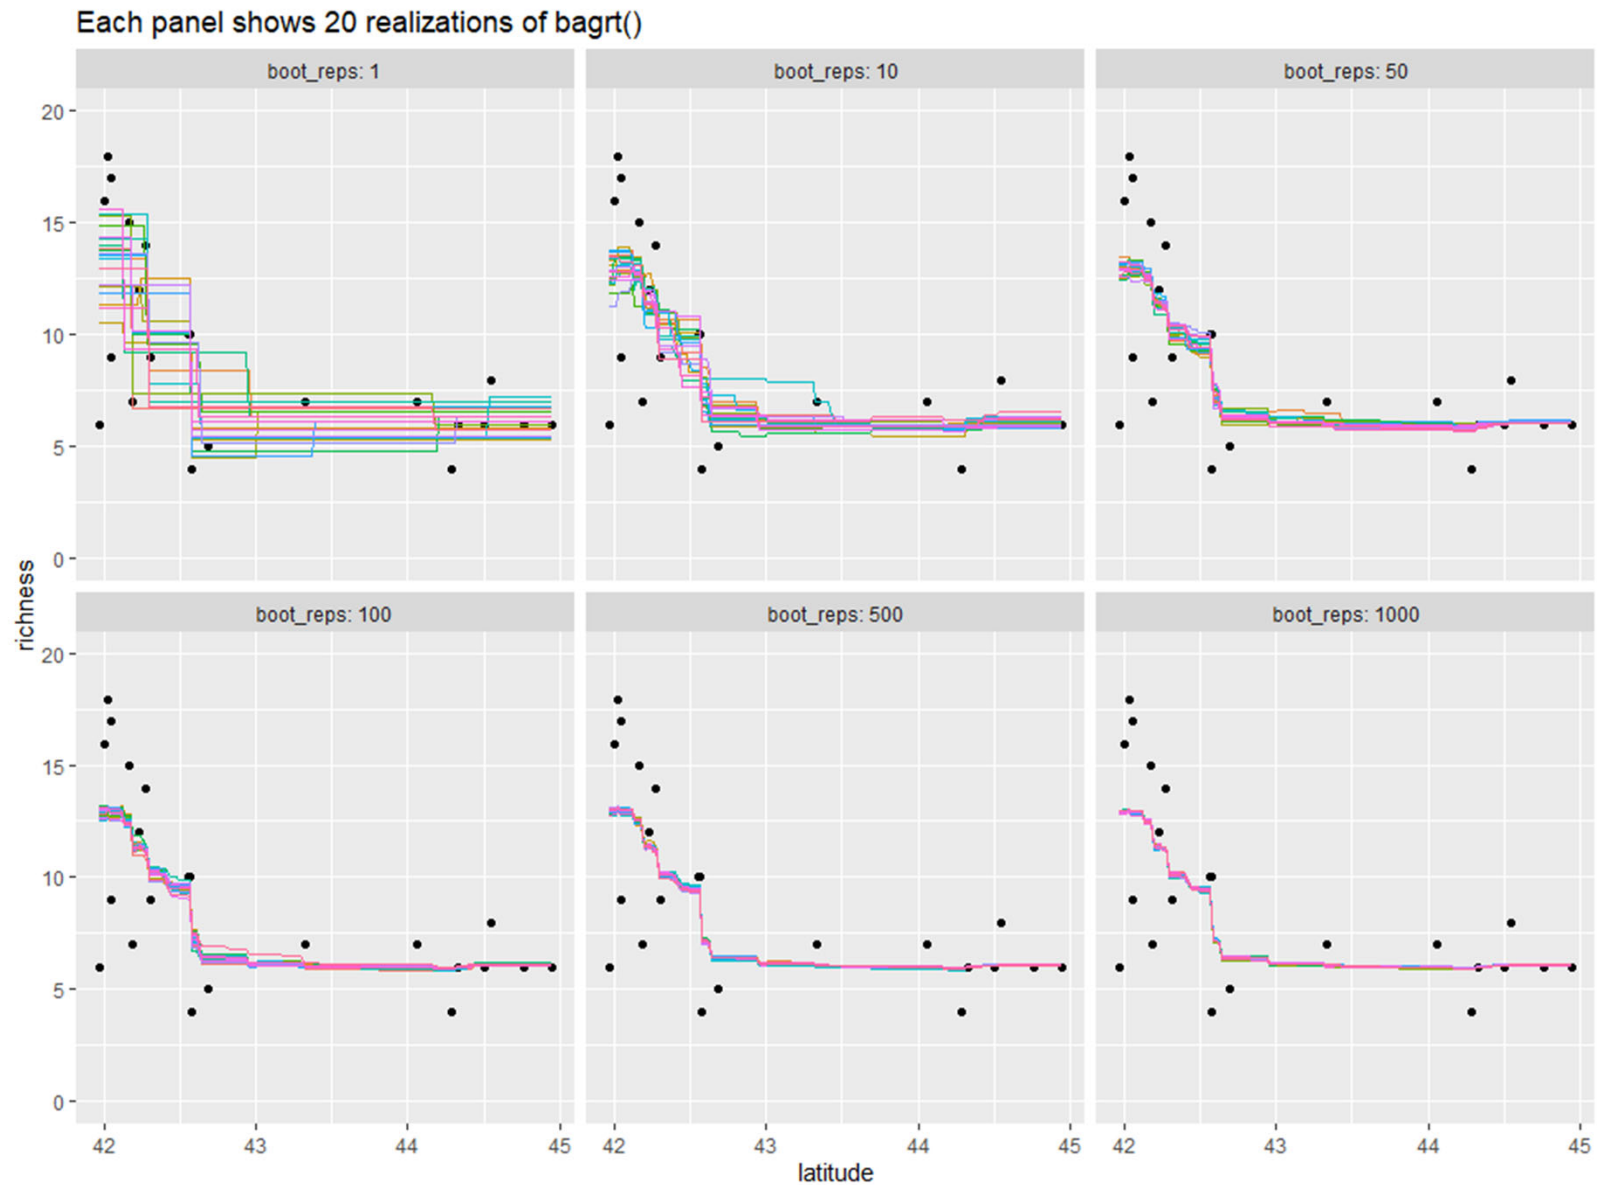
**Base model:** can be any type of model

Bagged regression tree (blue) vs single regression tree (black)

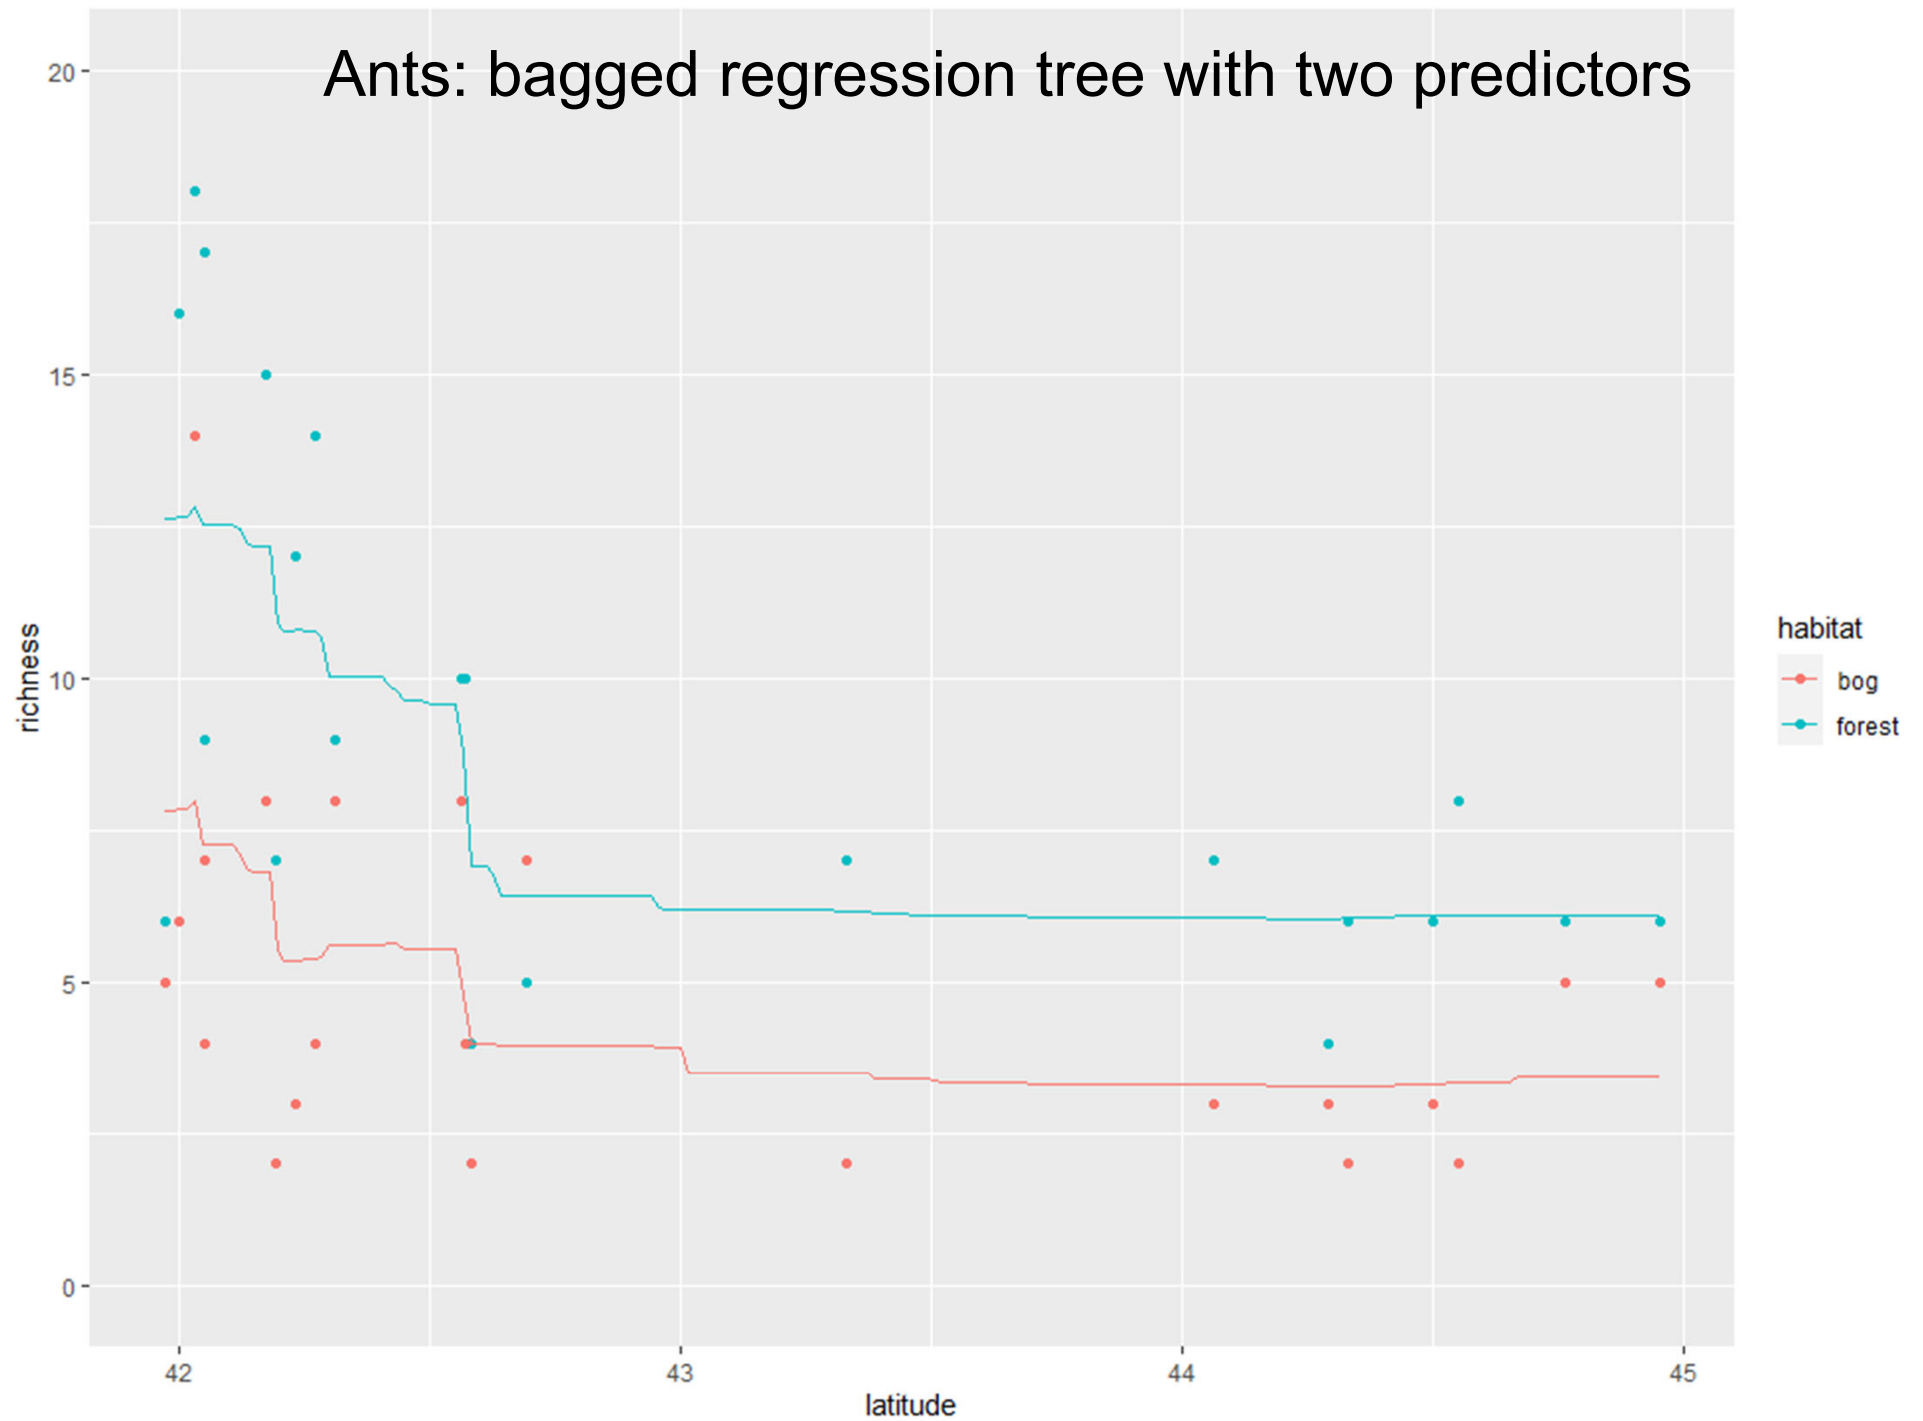




# Bagging reduces prediction variance



# Ants: bagged regression tree with two predictors



# Random forest

## Algorithm

for many repetitions

- randomly select  $m$  predictor variables

- resample the data (rows) with replacement


- train the tree model

- record prediction

final prediction = mean of predictions

# R packages

- randomForest
  - original Breimen (2001) algorithm
  - Fortran
- ranger
  - fast implementation
  - C++




# Boosting algorithm

- Too complex for this quick overview
- Basic idea: learn slowly by building up an ensemble iteratively from many models, each with small weight
- Key tuning parameter: learning rate
- Can include aspects of bagging and RF
- Training algorithm: gradient descent

# Boosting packages in R

- **gbm**: gradient boosting machines
  - boosted decision trees
  - good, stable, maintained
  - retired (no new features)
- **xgboost**: extreme gradient boosting
  - R interface to very fast C++ library
  - additional algorithm innovations
  - current industry standard





# Neural Networks in R

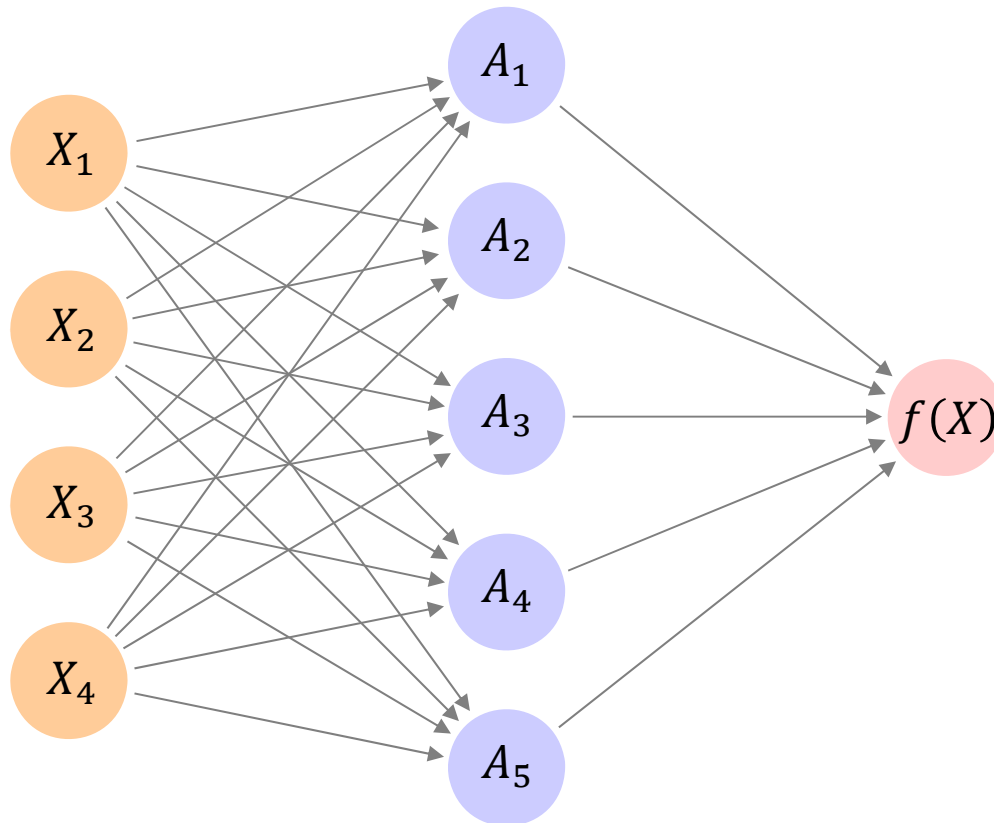
- keras
  - currently easiest to use
  - tensorflow (Google)
- torch
  - coming along (v 0.9)
  - PyTorch (Facebook, now Linux Foundation)

# Single layer NN

Input  
layer

Hidden  
layer

Output  
layer




# Single layer NN

Input  
layer

Hidden  
layer

Output  
layer



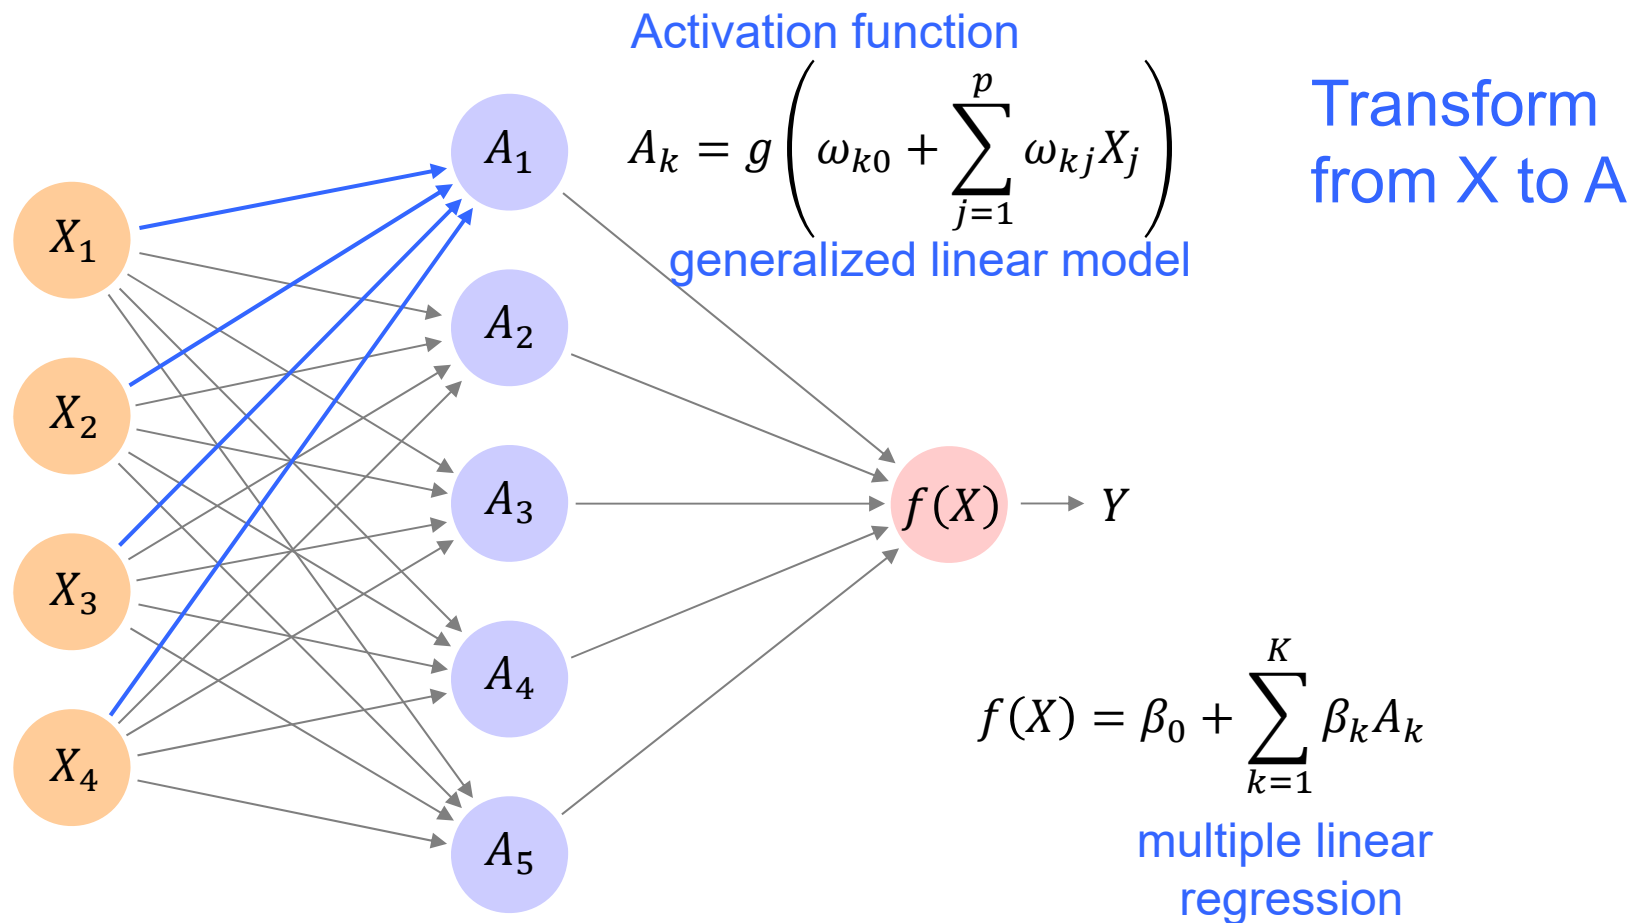
Transform  
from  $X$  to  $A$

# Single layer NN

Input  
layer

Hidden  
layer

Output  
layer




# Training algorithm

- Stochastic gradient descent
  - with back propagation


# Deep learning

- Multilayer neural networks
  - aka deep feedforward networks
  - aka multilayer perceptrons (MLP)
- Model algorithm
  - expressiveness
    - ability to approximate complex nonlinearity
  - architecture: width versus depth


Wide: 25 hidden units




# Deep: 25 hidden units









# Ants data: wide





# Ants data: deep

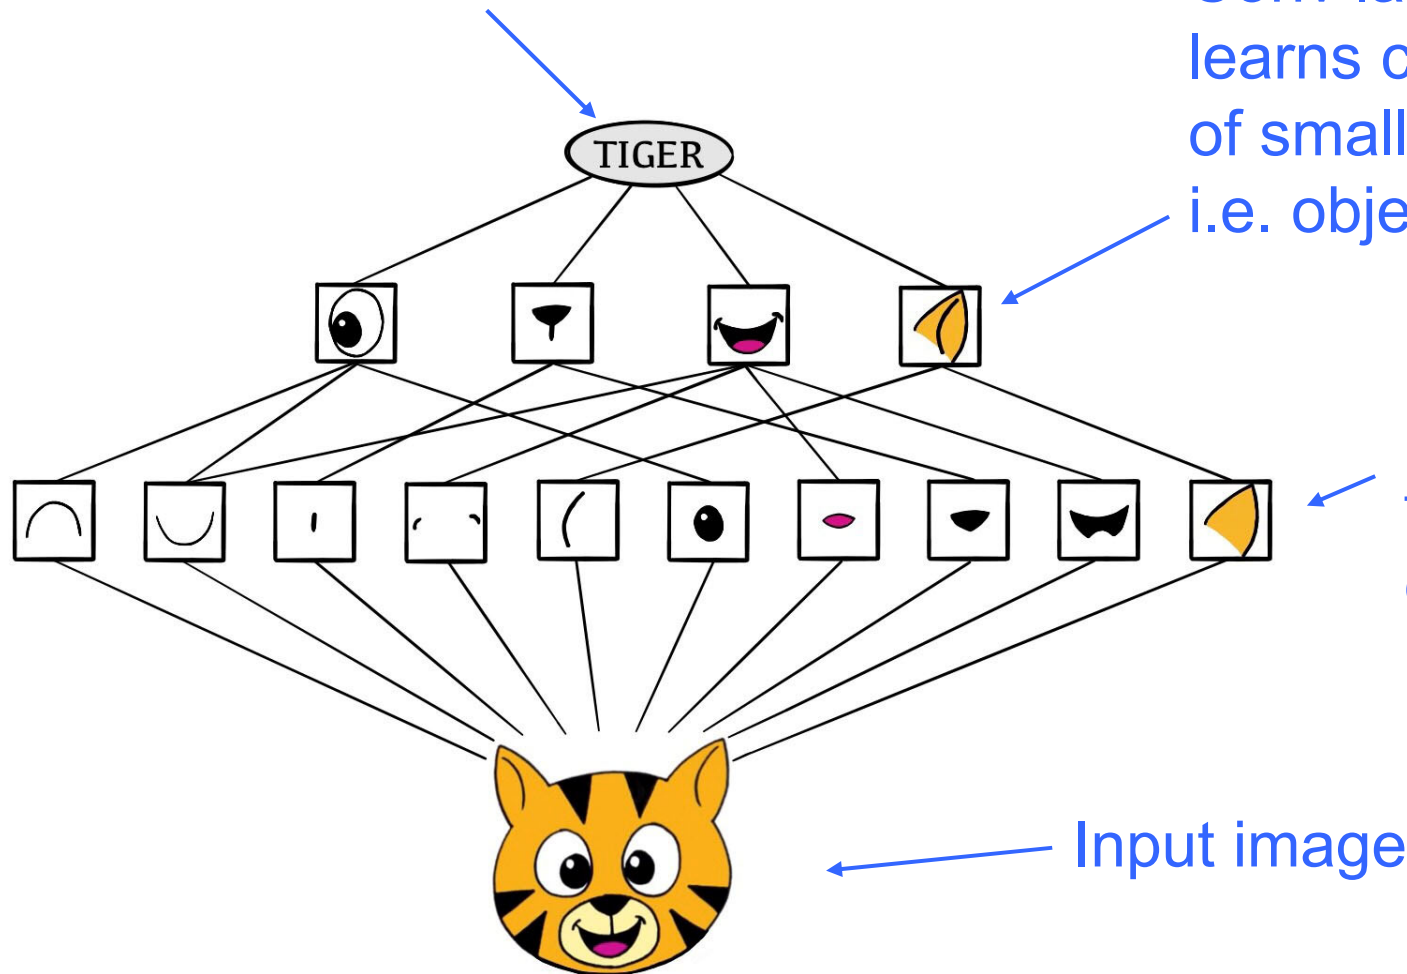


# Convolutional NNs

Output is high-level concept

Conv layer 2  
learns combinations  
of small features  
i.e. objects


Conv layer 1  
learns small  
features  
e.g. edges



# Transfer learning

Pretrained model on related big data

Retrain last 1-2 layers on specialized little data



# Rapid innovation

- Architectures
- Algorithms
- Recent example: transformer

# Outlook

Automated data collection  
+  
machine learning  
=  
revolutionary